

Diamond Price Prediction

Group :

- Mohamed Gamal
 - Amira Djaiz
 - Aya Abu Dabbur
 - Mai Serry
 - Dyaa Dwekat
-

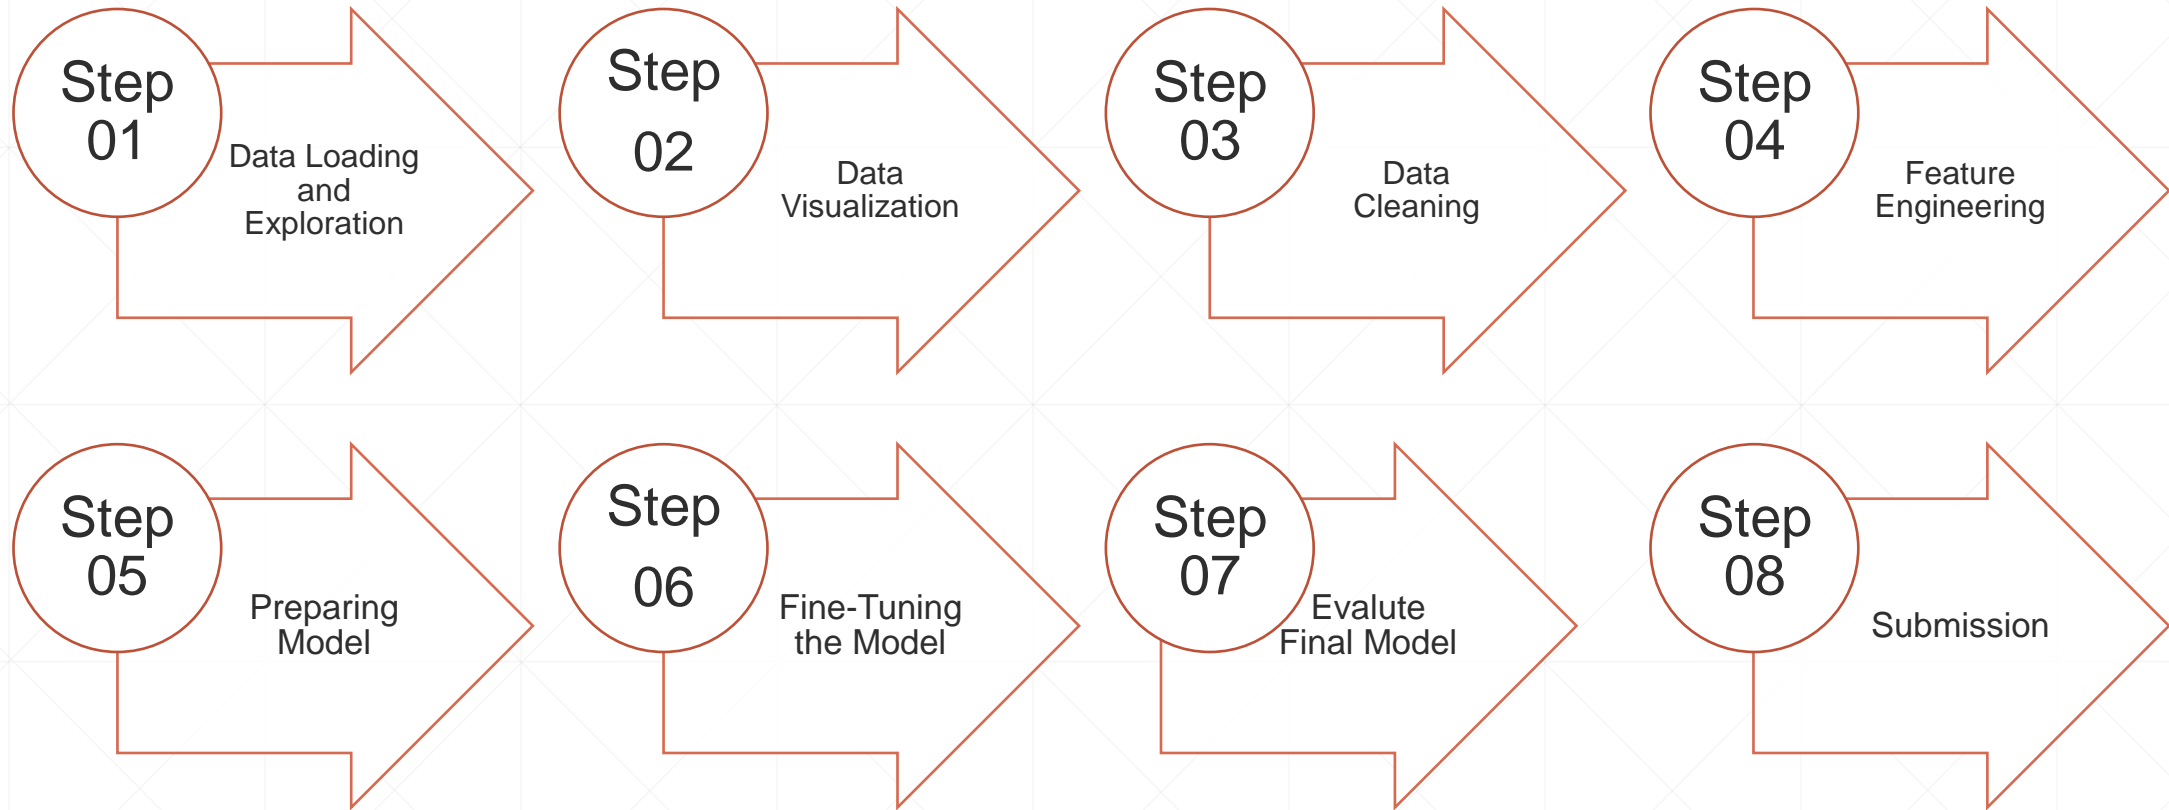
Supervisor:

- Saad Altamari
-

Content

- Introduction
- Data Loading and Exploration
- Data Visualization
- Data Cleaning
- Feature Engineering and Data Preparation
- Preparing Model
- Fine-Tuning the Model
- Submission

Structure



Introduction

Introduction

Background

- Understanding the complexities of diamond pricing, which depends on the 4Cs (Carat, Cut, Color, Clarity), is essential. These attributes interact in complex ways, making accurate price prediction a challenging task.

Objective

- Our goal is to develop a robust machine-learning model that can predict diamond prices accurately. This involves thorough data analysis and selecting the best predictive techniques.



Data Loading and Exploration

Data Loading

The first step of the project involved loading the training data.

```
df = pd.read_csv('/kaggle/input/diamond-price-predicton-2024/train.csv')
```

Data Exploration

`df.head()`

	Id	carat	cut	color	clarity	depth	table	price	x	y	z
0	1	1.06	Ideal	I	SI2	61.8	57.0	4270	6.57	6.60	4.07
1	2	1.51	Premium	G	VVS2	60.9	58.0	15164	7.38	7.42	4.51
2	3	0.32	Ideal	F	VS2	61.3	56.0	828	4.43	4.41	2.71
3	4	0.53	Ideal	G	VS2	61.2	56.0	1577	5.19	5.22	3.19
4	5	0.70	Premium	H	VVS2	61.0	57.0	2596	5.76	5.72	3.50

`df.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 43152 entries, 0 to 43151
Data columns (total 11 columns):
 #   Column      Non-Null Count  Dtype  
---  -
 0   Id          43152 non-null  int64  
 1   carat       43152 non-null  float64
 2   cut         43152 non-null  object  
 3   color       43152 non-null  object  
 4   clarity     43152 non-null  object  
 5   depth       43152 non-null  float64
 6   table       43152 non-null  float64
 7   price       43152 non-null  int64  
 8   x           43152 non-null  float64
 9   y           43152 non-null  float64
10  z           43152 non-null  float64
dtypes: float64(6), int64(2), object(3)
memory usage: 3.6+ MB
```


Data Preprocessing

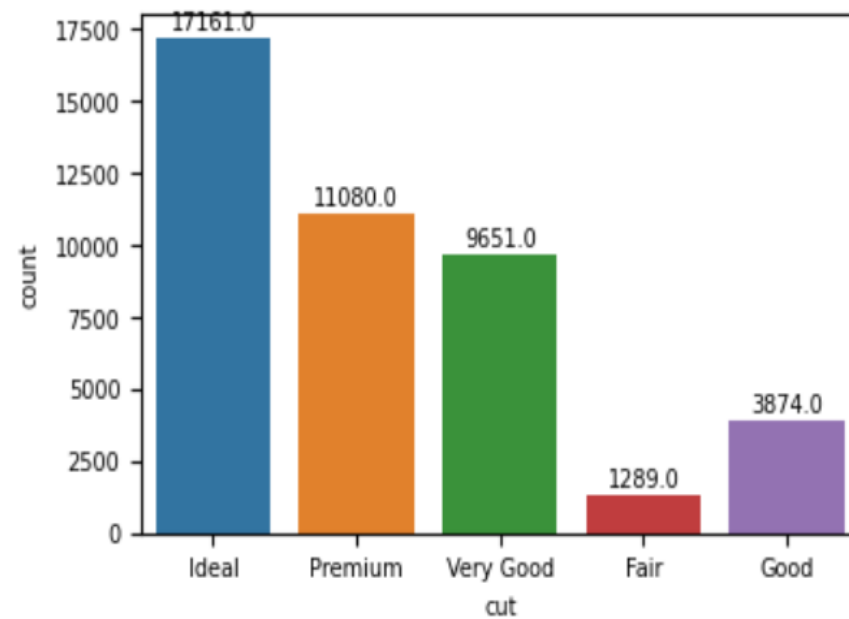
Removing Irrelevant Columns: The 'Id' column was determined to be unnecessary for our analysis and was thus removed from the dataset.

```
df_no_id = df.drop("Id", axis=1)
duplicates = df_no_id.duplicated()
num_duplicates = duplicates.sum()
num_duplicates
```

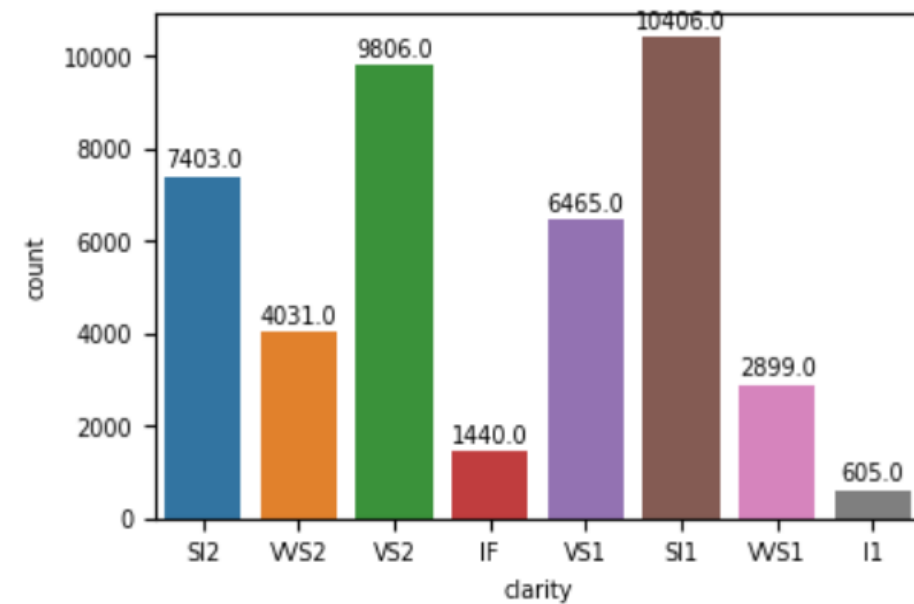
Data Visualization

Data Visualization

- Countplot for Cut




- Countplot for Clarity

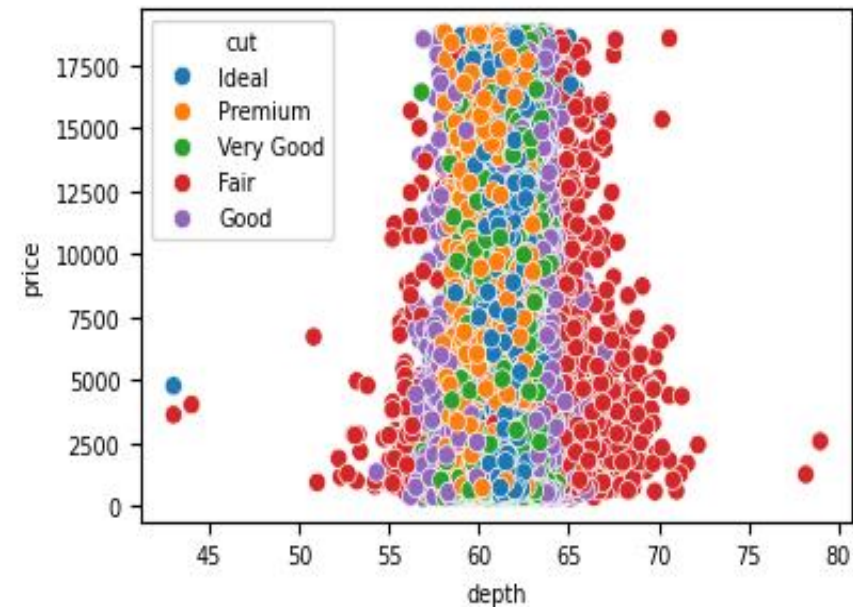


Data Visualization

- Scatterplot for Carat vs. Price:

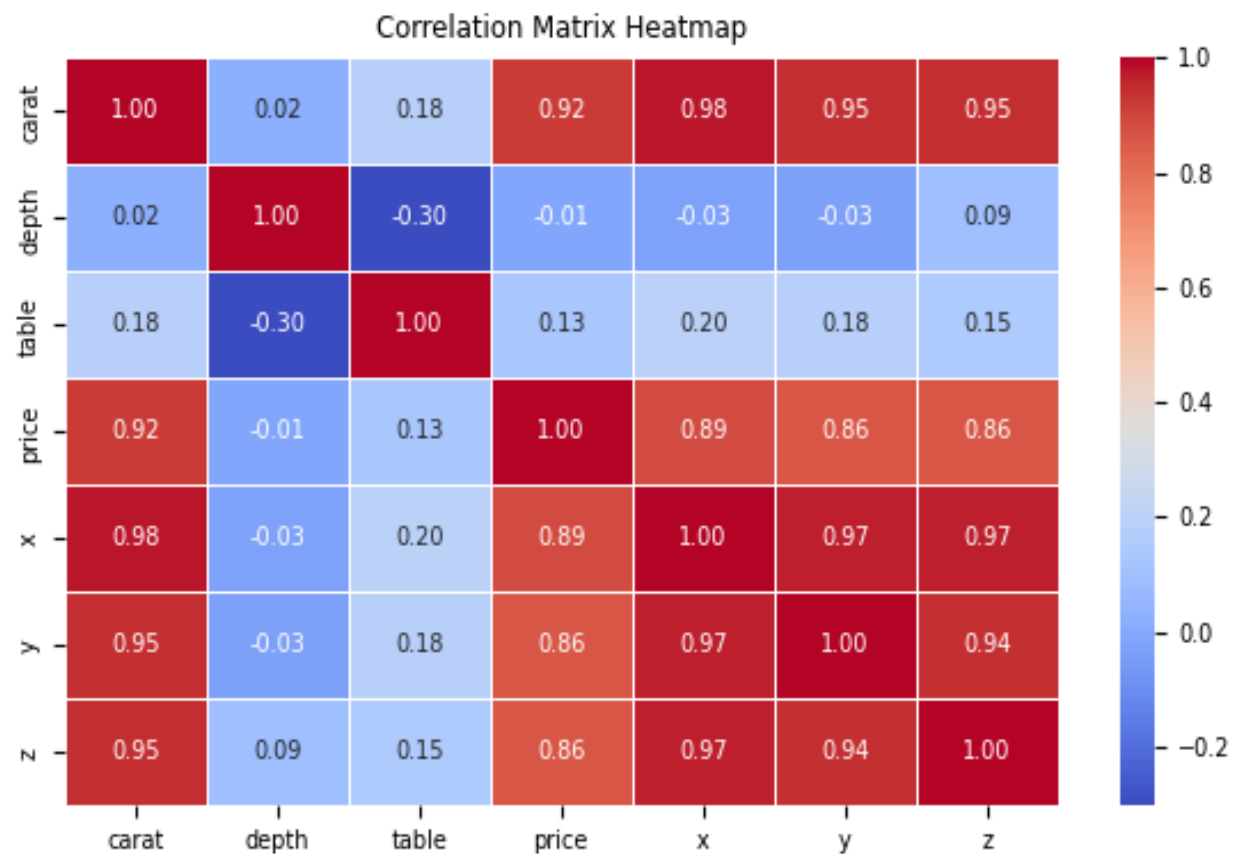


- Scatterplot for Depth vs. Price:



Data Visualization

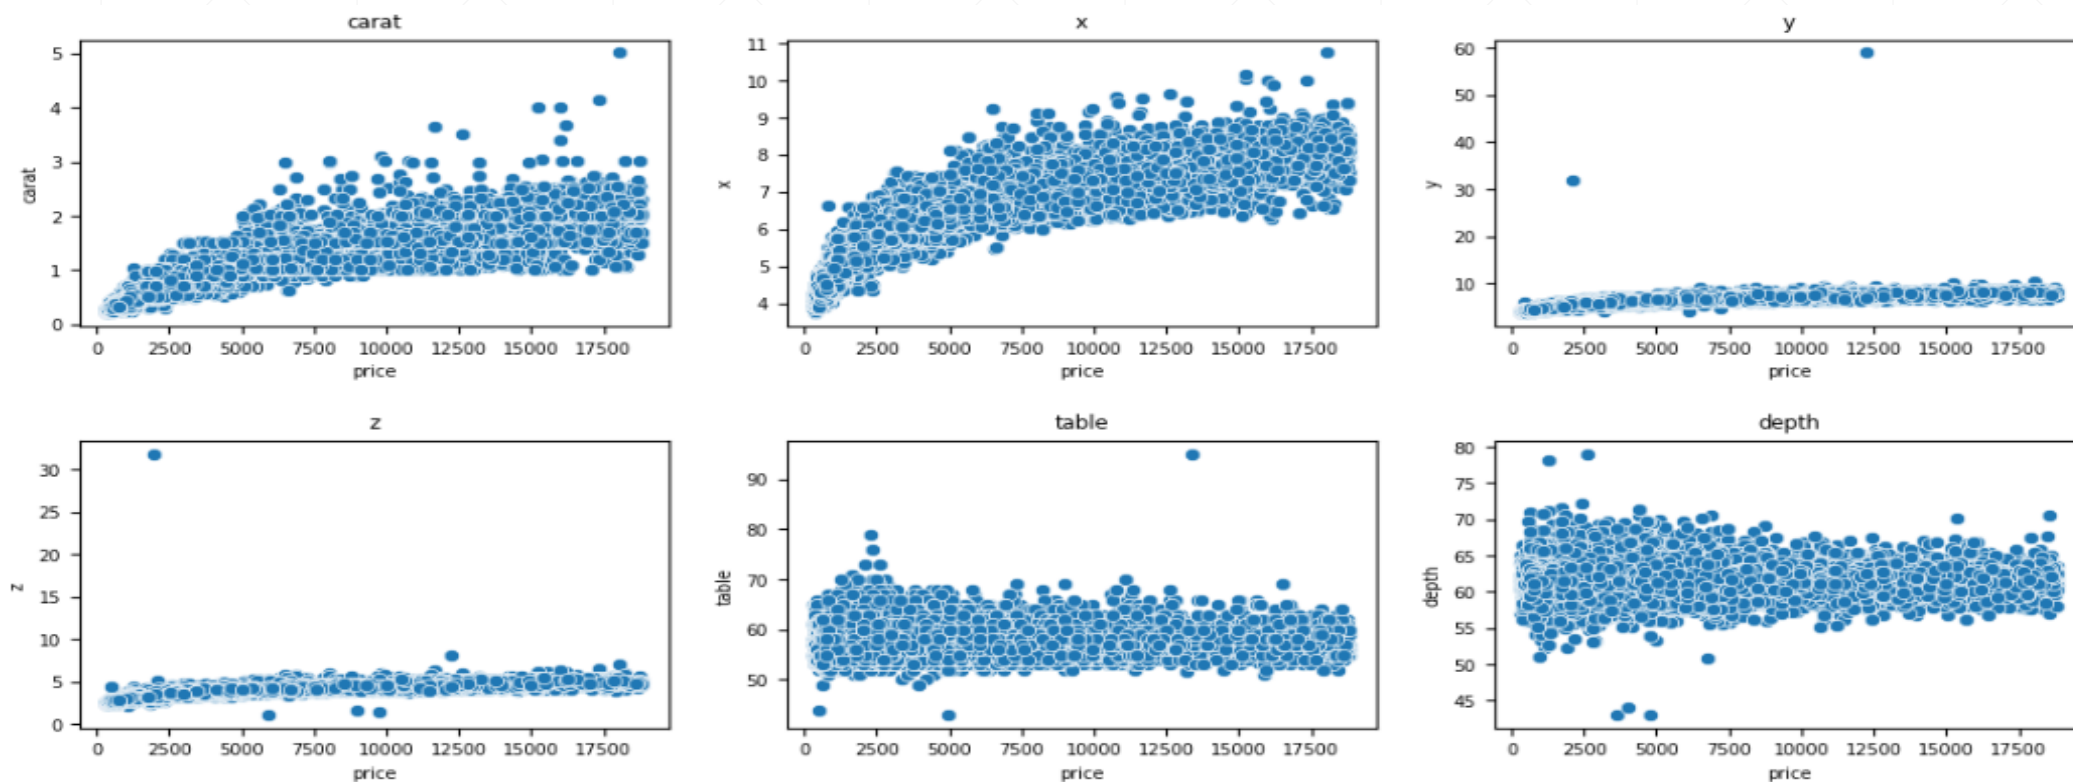
Heatmap for Feature Correlation (for numeric features)









Data Cleaning

Data Cleaning

Handling Outliers



Data Cleaning



Feature Engineering and Data Preparation

Feature Engineering



Feature Engineering

```
enc_processed_df.corr()
```


```
price      1.000000
carat      0.922247
volume     0.909367
table      0.127605
depth     -0.013130
cut        -0.055436
clarity    -0.146704
color      -0.170819
Name: price, dtype: float64
```

```
enc_processed_df.head()
```

	carat	cut	color	clarity	depth	table	price	volume
0	1.06	4.0	1.0	1.0	61.8	57.0	4270	740.430017
1	1.51	3.0	3.0	5.0	60.9	58.0	15164	957.036622
2	0.32	4.0	4.0	3.0	61.3	56.0	828	304.412538
3	0.53	4.0	3.0	3.0	61.2	56.0	1577	432.648888
4	0.70	3.0	2.0	5.0	61.0	57.0	2596	542.129615

Data Exploration

Correlation Matrix: Calculated the correlation matrix to examine the relationships between different features.



Preparing Models

Preparing Model

- **Custom Transformer for Feature Engineering**

We define a custom transformer `CombinedAttributesAdder` to generate a new feature, volume, based on the diamond dimensions:

```
from sklearn.base import BaseEstimator, TransformerMixin


x_ix, y_ix, z_ix, table_ix = [
    list(df.drop(["cut", "color", "clarity", "price"], axis=1).columns).index(col)
    for col in ("x", "y", "z", "table")]

class CombinedAttributesAdder(BaseEstimator, TransformerMixin):
```

Preparing Model

Splitting the Data

```
df.loc[:, "carat_cat"] = pd.cut(df["carat"],  
                                bins=[0.19, 0.5, 0.75, 1, 1.5, np.inf],  
                                labels=[1, 2, 3, 4, 5])  
  
df["carat_cat"].hist();
```



```
from sklearn.model_selection import StratifiedShuffleSplit  
  
split = StratifiedShuffleSplit(n_splits=10, test_size=0.2, random_state=42)  
  
for train_index, test_index in split.split(df, df["carat_cat"]):  
    train_set = df.iloc[train_index]  
    test_set = df.iloc[test_index]
```

Preparing Model

Numerical Pipeline

```
num_pipeline = Pipeline([
    ('imputer', SimpleImputer(strategy="median")),
    ('attrs_adder', CombinedAttributesAdder()),
    ('std_scaler', StandardScaler()),
])
```

Full Pipeline

```
full_pipeline = ColumnTransformer([
    ("num", num_pipeline, num_attribs),
    ("cat", OrdinalEncoder(categories=[colors_cats, cuts, clarities]), cat_attribs)
])
```

Data Transformation

```
x_train_prepared = full_pipeline.fit_transform(x_train)
x_test_prepared = full_pipeline.transform(x_test)
```



```
from sklearn.linear_model import LinearRegression
linear_model=LinearRegression()
linear_model.fit(x_train_prepared,y_train)
```

▼ LinearRegression
LinearRegression()

```
linear_model.score(x_train_prepared, y_train)
```

```
0.9186449950483981
```

Linear Regression model

We evaluate the model's performance on the training set.

```
from sklearn.metrics import mean_squared_error

lin_predic = linear_model.predict(x_train_prepared)
lin_rmse = mean_squared_error(
    y_train,
    lin_predic,
    squared=False #RMSE
)
lin_rmse
```

1133.7829679831411

Linear Regression model

Calculating RMSE on
Training Data.

```
from sklearn.model_selection import cross_val_score
```

```
lin_scores = cross_val_score(  
    linear_model,  
    x_train_prepared,  
    y_train,  
    scoring = "neg_mean_squared_error",  
    cv = 10  
)  
lin_rmse_scores = np.sqrt(-lin_scores)  
print("Scores: ", lin_rmse_scores)  
print("Mean: ", lin_rmse_scores.mean())  
print("Standard Deviation: ", lin_rmse_scores.std())
```

```
Scores: [1145.17612289 1137.04199626 1101.15552983 1208.70921887 1142.37399683  
1110.04639488 1143.66724006 1144.35294896 1098.31134199 1148.04160739]  
Mean: 1137.887639796304  
Standard Deviation: 29.93097691246276
```

Linear Regression model

use cross-validation to
evaluate the model's
performance.

```
from sklearn.tree import DecisionTreeRegressor
tree_model=DecisionTreeRegressor(random_state = 42)
tree_model.fit(x_train_prepared,y_train)
```

▼ DecisionTreeRegressor

```
DecisionTreeRegressor(random_state=42)
```

```
tree_model.score(x_train_prepared, y_train)
```

```
0.9999958224864528
```

Decision Tree Regressorion Model

We evaluate the model's
performance on the
training set.

```
tree_predic = tree_model.predict(x_train_prepared)
tree_rmse = mean_squared_error(
    y_train,
    tree_predic,
    squared=False
)
tree_rmse
```

8.142529435502995

Decision Tree Regressorion Model

Calculating RMSE on
Training Data.

Decision Tree Regression Model

use cross-validation to evaluate the model's performance.

```
tree_scores = cross_val_score(
    tree_model,
    x_train_prepared,
    y_train,
    scoring = "neg_mean_squared_error",
    cv = 10
)
tree_rmse_scores = np.sqrt(-tree_scores)
print("Scores: ", tree_rmse_scores)
print("Mean: ", tree_rmse_scores.mean())
print("Standard Deviation: ", tree_rmse_scores.std())
```

Scores: [742.84381883 703.8662325 764.20661659 697.67372586 772.21417935

757.46766405 754.37957042 720.71381911 695.87127516 730.35974988]

Mean: 733.9596651750455

Standard Deviation: 26.99672946198288

```
from sklearn.ensemble import RandomForestRegressor
```

```
forest_model = RandomForestRegressor(n_estimators=100, random_state=42)  
forest_model.fit(x_train_prepared, y_train)
```

```
▼ RandomForestRegressor  
RandomForestRegressor(random_state=42)
```

```
forest_model.score(x_train_prepared, y_train)
```

```
0.9972230352470094
```

RandomForestRegressor Model

We evaluate the model's performance on the training set.

```
forest_predic = forest_model.predict(x_train_prepared)
forest_rmse = mean_squared_error(
    y_train,
    forest_predic,
    squared=False
)
forest_rmse
```

204.63439946095136

RandomForest Regressor Model

Calculating RMSE on
Training Data.


```
forest_scores = cross_val_score(
    forest_model,
    x_train_prepared,
    y_train,
    scoring = "neg_mean_squared_error",
    cv = 10
)
forest_rmse_scores = np.sqrt(-forest_scores)
print("Scores: ", forest_rmse_scores)
print("Mean: ", forest_rmse_scores.mean())
print("Standard Deviation: ", forest_rmse_scores.std())
```

Scores: [563.51686332 536.22530208 575.45252782 536.89516107 558.59890765

563.49796513 558.90275627 519.22794338 527.196818 563.58505227]

Mean: 550.3099296989064

Standard Deviation: 17.82669031675242

RandomForest Regressor Model

use cross-validation to
evaluate the model's
performance.

Fine-Tuning the Model

```

from sklearn.model_selection import RandomizedSearchCV

param_distributions = {
    'n_estimators': np.random.randint(1, 200, 10),
    'max_features': np.random.randint(1, 8, 10),
}

forest_model_fine = RandomForestRegressor(random_state=42)
rnd_search = RandomizedSearchCV(
    forest_model_fine,
    param_distributions=param_distributions,
    n_iter=10,
    cv=5,
    scoring='neg_mean_squared_error',
    random_state=42
)
rnd_search.fit(x_train_prepared, y_train)

```

```

RandomizedSearchCV(cv=5, estimator=RandomForestRegressor(random_state=42),
,
    param_distributions={'max_features': array([7, 7, 2, 4, 1, 2, 1, 4, 1, 2]),
,
    'n_estimators': array([ 72, 108, 92, 24, 68, 1, 22, 46, 62, 85])},
,
    random_state=42, scoring='neg_mean_squared_error')

```

Fine-Tuning the Model

```
cvres = rnd_search.cv_results_  
for mean_score, params in zip(cvres["mean_test_score"], cvres["params"]):  
    print(np.sqrt(-mean_score), params)
```

```
666.0574858358463 {'n_estimators': 24, 'max_features': 1}  
581.8702387052098 {'n_estimators': 24, 'max_features': 2}  
548.3524475625106 {'n_estimators': 72, 'max_features': 4}  
1059.9121520738188 {'n_estimators': 1, 'max_features': 1}  
644.1725377061641 {'n_estimators': 68, 'max_features': 1}  
547.5641124850662 {'n_estimators': 85, 'max_features': 4}  
564.4709532633339 {'n_estimators': 92, 'max_features': 2}  
643.6625766840698 {'n_estimators': 72, 'max_features': 1}  
555.1399398111733 {'n_estimators': 72, 'max_features': 7}  
555.1399398111733 {'n_estimators': 72, 'max_features': 7}
```

```
final_model = rnd_search.best_estimator_  
final_model
```

```
RandomForestRegressor(max_features=4, n_estimators=85, random_state=42)
```

Fine-Tuning the Model

The best-performing set of parameters is {'n_estimators': 72, 'max_features': 4} which resulted in the lowest RMSE of 548.35.

Evaluate Final Model

```
# final_predictions = final_model.predict(x_test_prepared)
final_predictions = final_model.predict(x_train_prepared)

# final_mse = mean_squared_error(y_test, final_predictions)
final_mse = mean_squared_error(y_train, final_predictions)
final_rmse = np.sqrt(final_mse)
final_rmse
```

202.77888794641905

```
from scipy import stats

confidence = 0.95
squared_errors = (final_predictions - y_train) ** 2
np.sqrt(stats.t.interval(confidence, len(squared_errors) - 1,
                          loc=squared_errors.mean(),
                          scale=stats.sem(squared_errors)))
```

array([198.18512292, 207.27086589])

Evaluate Final Model

Calculate the confidence interval for the Root Mean Square Error (RMSE) of the final model's predictions.

Submission

Submission

the first few rows of our submission file:

```
data={'Id':Id,'price':price}
sub=pd.DataFrame(data)
filename=f'mg_sub_{int(round(final_rmse))}.csv'
print(filename)
print(sub)
sub.to_csv(filename, index=False)
```

Id	price
1	839.1597
2	2905.8908
3	875.2437
4	2753.8571
5	1054.3025
...	...
10786	4347.8571
10787	4786.2605
10788	13905.1765